Liam Dugan

ldugan@seas.upenn.edu | linkedin.com/in/liam-dugan | liamdugan.com

Research Focus

My research focuses on human and automated detection of AI-generated content. In particular, I am interested in the technical limitations and societal ramifications of AI detection tools and how we might deploy accurate AI detectors with minimal harm. More broadly I am interested in developing a deep understanding of LLMs: their stylistic tendencies, their internal representational dynamics, and their reasoning capabilities. My strengths are generally in large-scale data analysis, model inference, and software engineering.

Education

University of Pennsylvania — Philadelphia, PA

Ph.D, Computer Science (Advisor: Chris Callison-Burch) — *Aug. 2021 – est. May 2026*

M.S.E, Robotics — *Aug. 2017 – Dec. 2020*

B.S.E, Computer Engineering & East Asian Studies — *Aug. 2015 – Aug. 2020*

Work Experience

Summer 2025 **Google DeepMind** (New York, NY) - *Student Researcher*
Hosts: Philip Pham & Matthew Denton
Project: *Mechanistic Interpretability for Faithful & Explainable AI Detection*

Summer 2022 **Roblox** (San Mateo, CA) - *PhD Research Intern*
Hosts: Morgan McGuire & Victor Zordan
Project: *Real-Time Speech-to-Speech Translation*

Summer 2021 **John's Hopkins University** (Baltimore, MD) - *Visiting Research Scholar*
Hosts: Kevin Duh, Paul McNamee, Matt Post
Project: *Machine Translation for Cross-Lingual Information Retrieval*

Summer 2019 **NVIDIA** (Santa Clara, CA) - *Autonomous Driving Software Intern*
Host: Gajanan Bhat
Project: *Docker Image Server for Autonomous Driving*

Summer 2018 **Forterra** (Clarksburg, MD) - *Software Engineering Intern*
Host: Anne Schneider
Project: *Velodyne VLP-16 LIDAR Point Cloud Classifiers*

Publications

Preprints

2025 Ryuto Koike*, **Liam Dugan***, Masahiro Kaneko, Chris Callison-Burch, and Naoaki Okazaki. Machine-Generated Text Detectors are Membership Inference Attacks. ArXiv, October 2025

Meiqing Jin*, **Liam Dugan***, and Chris Callison-Burch. Controlling Difficulty of Generated Text for AI-Assisted Language Learning. ArXiv, June 2025

Arihant Tripathi*, **Liam Dugan***, Charis Gao, Maggie Huan, Meiqing Jin, Peter Zhang, David Zhang, Julia Zhao, and Chris Callison-Burch. Domain Gating Networks for AI-Generated Text Detection. ArXiv, May 2025

Minseok Jung, Cynthia Fuertes Panizo, **Liam Dugan**, May Fung, Pin-Yu Chen, and Paul Pu Liang. Group-Adaptive Threshold Optimization for Robust AI-Generated Text Detection. ArXiv, February 2025

2023 Josh Ludan, Qing Lyu, Yue Yang, **Liam Dugan**, Mark Yatskar, and Chris Callison-Burch. Interpretable-by-Design Text Understanding with Iteratively Generated Concept Bottleneck. ArXiv, October 2023

Conference Papers

2025 **Liam Dugan**, Andrew Zhu, Firoj Alam, Preslav Nakov, Marianna Apidianaki, and Chris Callison-Burch. GenAI Content Detection Task 3: Cross-Domain Machine Generated Text Detection Challenge. In *Proceedings of the 1st Workshop on GenAI Content Detection (GenAIDetect)*, pages 377–388, Abu Dhabi, UAE, January 2025

2024 Andrew Zhu, **Liam Dugan**, and Chris Callison-Burch. ReDel: A toolkit for LLM-powered recursive multi-agent systems. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 162–171, Miami, Florida, USA, November 2024. Association for Computational Linguistics

Runsheng Huang, **Liam Dugan**, Yue Yang, and Chris Callison-Burch. MiRAGeNews: Multimodal realistic AI-generated news detection. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 16436–16448, Miami, Florida, USA, November 2024. Association for Computational Linguistics

Andrew Zhu, Alyssa Hwang, **Liam Dugan**, and Chris Callison-Burch. FanOutQA: A Multi-Hop, Multi-Document Question Answering Benchmark for Large Language Models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 18–37, Bangkok, Thailand, August 2024

Liam Dugan, Alyssa Hwang, Filip Trhlík, Andrew Zhu, Josh Magnus Ludan, Hainiu Xu, Daphne Ippolito, and Chris Callison-Burch. RAID: A Shared Benchmark for Robust Evaluation of Machine-Generated Text Detectors. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 12463–12492, Bangkok, Thailand, August 2024. **Nominated for Outstanding Paper Award**

2023 Andrew Zhu*, **Liam Dugan***, Alyssa Hwang, and Chris Callison-Burch. Kani: A Lightweight and Highly Hackable Framework for Building Language Model Applications. In *Proceedings of the 3rd Workshop for Natural Language Processing Open Source Software (NLP-OSS 2023)*, pages 65–77, Singapore, Singapore, December 2023. Empirical Methods in Natural Language Processing

Josh Ludan, Qing Lyu, Yue Yang, **Liam Dugan**, Mark Yatskar, and Chris Callison-Burch. Interpretable-by-Design Text Understanding with Iteratively Generated Concept Bottleneck. ArXiv, October 2023

Liam Dugan, Anshul Wadhawan, Kyle Spence, Chris Callison-Burch, Morgan McGuire, and Victor Zordan. Learning When to Speak: Latency and Quality Trade-offs for Simultaneous Speech-to-Speech Translation with Offline Models. In *Proc. INTERSPEECH 2023*, pages 5265–5266, August 2023

Hannah Gonzalez, **Liam Dugan**, Eleni Miltsakaki, Zhiqi Cui, Jiaxuan Ren, Bryan Li, Shriyash Upadhyay, Etan Ginsberg, and Chris Callison-Burch. Enhancing Human Summaries for Question-Answer Generation in Education. In *Proceedings of the 18th Workshop on Innovative Use of NLP for Building Educational Applications (BEA 2023)*, pages 108–118, Toronto, Canada, July 2023. Association for Computational Linguistics

Li Zhang*, **Liam Dugan***, Hainiu Xu*, and Chris Callison-Burch. Exploring the Curious Case of Code Prompts. In *Proceedings of the 1st Workshop on Natural Language Reasoning and Structured Explanations (NLRSE)*, pages 9–17, Toronto, Canada, June 2023. Association for Computational Linguistics **Selected for Oral Presentation**

Liam Dugan*, Daphne Ippolito*, Arun Kirubarajan, Sherry Shi, Chris Callison-Burch. Real or Fake Text?: Investigating Human Ability to Detect Boundaries between Human-Written and Machine-Generated Text. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37(11), pages 12763–12771, Washington, D.C., June 2023. **Selected for Oral Presentation**

Aarohi Srivastava, Abhinav Rastogi, and **(440 others)**. Beyond the Imitation Game: Quantifying and extrapolating the capabilities of language models. *Transactions on Machine Learning Research*, May 2023

2022 Daphne Ippolito, **Liam Dugan**, Emily Reif, Ann Yuan, Andy Coenen, and Chris Callison-Burch. The Case for a Single Model that can Both Generate Continuations and Fill-in-the-Blank. In *Findings of the Association for Computational Linguistics: NAACL 2022*, pages 2421–2432, Seattle, United States, July 2022

Liam Dugan, Eleni Miltsakaki, Shriyash Upadhyay, Etan Ginsberg, Hannah Gonzalez, DaHyeon Choi, Chuning Yuan, and Chris Callison-Burch. A Feasibility Study of Answer-Agnostic Question Generation for Education. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 1919–1926, Dublin, Ireland, May 2022

2020 **Liam Dugan***, Daphne Ippolito*, Arun Kirubarajan*, and Chris Callison-Burch. RoFT: A Tool for Evaluating Human Detection of Machine-Generated Text. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 189–196, Online, October 2020. Association for Computational Linguistics

2019 Zhengyi Luo, Austin Small, **Liam Dugan** and Stephen Lane. Cloud Chaser: Real Time Deep Learning Computer Vision on Low Computing Power Devices. In *Eleventh International Conference on Machine Vision (ICMV 2018)*, volume 11041, page 110412Q. International Society for Optics and Photonics, SPIE, 2019

Invited Talks

2025 Detecting AI-Generated Content in the Real World: Drexel University, Philadelphia, May 2025 (slides)

Progress Towards Robust and Deployable AI Detectors in the Real World: Stanford University Hoover Institute, Palo Alto CA, February 2025 (slides)

Stylistic Signatures of LLMs and How to Detect Them: University of Pennsylvania ILST Seminar, Philadelphia PA, February 2025 (slides) (recording)

2024 RAID: A Shared Benchmark for Robust Evaluation of Machine-Generated Text Detectors: University of Maryland, College Park MD, June 2024 (slides)

2023 Should we still use Text for Speech-to-Speech Translation? Promise meets Practice: John's Hopkins University, Baltimore MD, May 2023 (slides)

Real or Fake Text: Investigating Human Ability to Detect Boundaries between Human-Written and Machine Generated Text: Brown University, Providence RI, March 2023 (slides)

Detecting Generated Text from ChatGPT and other LLMs: University of Pennsylvania, Philadelphia PA, February 2023 (slides)

2022 Intro to Machine Learning and AI Research: St. Joseph's Preparatory High School, Philadelphia PA, Feb. 2022 (slides)

Are Humans Able to Detect Boundaries between Human-Written and Machine-Generated Text?: University of Pennsylvania Computational Linguistics Lunch (CLUNCH), Philadelphia PA, Jan. 2022

News Articles

(11/28/25) Can teachers spot AI writing? Penn researchers weigh in - WHYY
(11/14/25) How chatbots make us feel - The Pulse (Podcast Interview)
(4/30/25) When ChatGPT Broke an Entire Field: An Oral History - Quanta Magazine
(9/9/24) AI Detectors are Easily Fooled, Researchers Find - EdScoop
(9/3/24) Teachers still can't trust AI text checkers - Axios
(8/22/24) Most AI text detectors aren't as reliable as advertised, study finds - TechBrew
(8/14/24) Putting AI Text Detectors to the Test: From Hype to Hard Data - SafeAI@Penn Newsletter
(8/12/24) Detecting Machine-Generated Text: An Arms Race With the Advancements of Large Language Models - Penn Engineering Today
(5/21/24) Originality.ai is the Most Accurate AI Detector According to an Extensive Study "RAID" - Originality.ai Blog
(1/2/24) Can Humans Learn To Spot Fake Text? - Penn Engineering Magazine
(9/20/23) Unlocking AI Potential: Unveiling Kani, the Groundbreaking Open-Source Framework Revolutionizing Large Language Model Applications - CJ&CO
(9/19/23) Researchers from the University of Pennsylvania Introduce Kani: A Lightweight, Flexible, and Model-Agnostic Open-Source AI Framework for Building Language Model Applications - MarkTechPost
(9/19/23) Kani: A Lightweight and Customizable Framework for Language Model Applications - TS2
(8/7/23) AI 'Watermarking' Tools Emerging to Tag Machine-Made Content - Bloomberg Law
(7/27/23) CNN Features Penn Engineering AI Research - Penn Engineering Today
(7/19/23) Academic Integrity and AI: Is Detection the Answer? - Temple Center for Teaching
(7/11/23) Bot or not? How to tell when you're reading something written by AI - CNN
(5/18/23) NewsChannel12 Investigates: Artificial Intelligence Part III - ABC News North Carolina
(4/26/23) Alien Minds, Immaculate Bullshit, Outstanding Questions - The Pennsylvania Gazette
(4/18/23) How can people navigate AI-generated misinformation? - Canvas 8
(4/11/23) Reddit Moderators Brace for a ChatGPT Spam Apocalypse - Vice
(3/10/23) Real or fake text? We can learn to spot the difference - Penn Today
(3/8/23) A Bot Isn't Going to Take Your Place, But AI Will Make Your Job Harder - Corporate Compliance Insights
(3/8/23) New Study Shows People Can Learn to Spot Machine-Generated Text - UniteAI
(3/6/23) How can humans detect AI writing? These Penn researchers have some tips - Technically Philly

(3/3/23) Can Humans Detect Text by AI Chatbot GPT? - Psychology Today
(3/2/23) People can learn to detect AI writing - Cosmos Magazine
(2/27/23) Real or Fake Text? We Can Learn to Spot the Difference - Penn Engineering Today
(12/19/22) How to spot AI-generated Text - MIT Technology Review
(1/23/18) Object-Seeking Robot Wins PennApps XVII - Penn Engineering Today
(9/10/17) At PennApps XVI, students made inter-dimensional robots and hung out with the founder of Quora - The Daily Pennsylvanian

Teaching

Summer 2023 **Teaching Assistant for CIS530, *Computational Linguistics***
Taught by Chris Callison-Burch. Wrote homework "Fine-Tuning Pre-Trained Language Models"

Fall 2022 **Teaching Assistant for CIS700, *Research Practicum***
Taught by Chris Callison-Burch

Spring 2022 **Teaching Assistant for CIS700, *Interactive Fiction & Text Generation***
Co-Taught by Chris Callison-Burch and Lara Martin

Fall 2021 **Teaching Assistant for CIS565, *GPU Programming & Architecture***
Taught by Shehzan Mohammed. Gave two guest lectures, "Optimizing Machine Learning with CUDA" and "Introduction to Machine Learning". Mentored students with ML final projects

Fall 2020 **Teaching Assistant for CIS530, *Computational Linguistics***
Taught by Clayton Greenberg. Wrote homework "Transformers and State-of-the-Art Models"

Spring 2020 **Teaching Assistant for CIS530, *Computational Linguistics***
Taught by Chris Callison-Burch. Wrote homework "Neural Machine Translation"

Fall 2019 **Head Teaching Assistant for CIS380, *Operating Systems***
Taught by Boon Thau Loo. Re-wrote homework write-ups and developed autograders
Gave guest lecture "Linux Page Replacement Algorithms and Belady's Anomaly"
Achieved highest ever course rating in TA Quality (3.37/4), and Overall Quality (3.29/4)

Spring 2019 **Teaching Assistant for CIS548, *Operating Systems***
Taught by Boon Thau Loo.

Fall 2018 **Teaching Assistant for CIS380, *Operating Systems***
Taught by Boon Thau Loo.

Spring 2018 **Teaching Assistant for CIS240, *Intro to Computer Systems***
Taught by Thomas Farmer.

Fall 2017 **Teaching Assistant for CIS240, *Intro to Computer Systems***
Taught by Camillo Jose Taylor.

Fall 2017 **Teaching Assistant for SD4x, *Programming for the Web with JavaScript***
Co-Taught by Chris Murphy and Swapneel Sheth.

Spring 2017 **Teaching Assistant for CIS240, *Intro to Computer Systems***
Taught by Thomas Farmer.

Fellowships, Awards, And Honors

Nov 2024 **Outstanding Reviewer EMNLP 2024**
Award given in recognition of my efforts when reviewing papers for the EMNLP 2024 conference

Aug 2022 **Roblox Research Grant**
Funding granted to continue work into speech-to-speech translation for the 2022-2023 academic year

Oct 2021 **Google Cloud Platform Research Grant**
For the server and compute costs of the Real or Fake Text website (http://roft.io)

May 2020 **Penn Engineering Exceptional Service Award**
For my work as Head Teaching Assistant for CIS380 (Operating Systems)

May 2019 **Penn Engineering Senior Design Award**
For my Senior Design Project (Scene++) [see video]

Oct 2018 **Foreign Language and Area Studies Undergraduate Fellowship**
Funding granted to continue my Master's Thesis research into east asian language NLP applications

Jan 2018 **Grand Prize & Best use of Cloud Hosting: PennApps XVII**
For my project Cloud Chaser (1st place out of 156 teams) [see video]

Sept 2017 **Third Prize: PennApps XVI**
For my project Todd: The Inter-Dimensional Robot (3rd place out of 158 teams) [see video]

Mentorship

2025 Arihant Tripathi, Maggie Huan, Charis Gao, David Zhang, Julia Zhao, Peter Zhang

—Project: Domain Gating Networks for AI Detection **[Submitted to AAAI 2026]**

Meiqing Jin (Independent Study)

—Project: Beginner-Friendly AI language tutors **[Submitted to EACL 2026]**

Amay Tripathi, Hongshuo Zhou, Andre van de Ven, Vignesh Lakshmanan

—Project: Tokenization-Free Neural Linguistic Steganography

2024 Tony An, Andrew Jiang, Ishaan Lal, Joseph Lee, Nathaniel Lao (Senior Design)

—Project: En Poisson **[Won 1st Place in CS Senior Design]**

Josh Magnus Ludan (Independent Study) — Current Position: PhD at University of Pennsylvania

Runsheng (Anson) Huang (Independent Study)

Filip Trhlik (Independent Study)

2023 Maya Guru, Yiran Chen, Sahit Penmatcha, Kaitlynn Soo, V. Veeramachaneni (Senior Design)

—Project: Dubble **[Won M&T Integration Lab Finalist & Judge Harold Berger Award]**

River Yijiang Dong (Independent Study) — Current Position: PhD at Cambridge University

Hainiu Xu (Independent Study) — Current Position: PhD at King's College London

Hannah Gonzalez (Independent Study) — Current Position: PhD at John's Hopkins University

Charlie Chen (Independent Study)

Anshul Wadhawan (Master's Thesis)

2022 Shriyash Upadhyay & Etan Ginsberg (Independent Study)

—Co-Founders at Martian **[Raised $9M Seed Funding]**

Service

Reviewing: EACL 2026, AAAI 2026, EMNLP 2025, COLING 2025, EMNLP 2024, ACL 2024, ACL 2023, ACL 2021
Organization: GenAI Detection Workshop @ COLING 2025, CLunch Fall 2024, PennNLP Reading Group 2024-Present
Conference Attendance: COLING 2025, NeurIPS 2024, EMNLP 2024, COLM 2024, ACL 2024, EMNLP 2023, Interspeech 2023, ACL 2023, AAAI 2023, NAACL 2022, ACL 2022

Technical Skills

Natural Languages: English (native), Japanese (advanced, business fluent - 7+ years [JLPT N2])
Programming Languages: Python, C, C++, Java, bash, CUDA, MATLAB, JavaScript, HTML/CSS, Verilog, Go